

# CONCURRENCY: SEQUENTIAL CONSISTENCY, DATA RACES, AND DYNAMIC ANALYSES

---

Lecture by Rohan Padhye

17-355/17-665/17-819: Program Analysis

Material from past lectures by Jonathan Aldrich, based in large part on slides by John Erickson, Stephen Freund, Madan Musuvathi, Mike Bond, and Man Cao

# Lecture Goals

- What is sequential consistency and why is it important?
- What is a data race, and what is data-race-free execution?
- Subtleties of data races and memory models
  - Why taking advantage of “harmless races” is almost certainly a bad idea
- Lockset analysis for data race detection
- Happens-before based data race detection
  - And high performance implementations, e.g. as in FastTrack

# SEQUENTIAL CONSISTENCY

---

# First things First

## Assigning Semantics to Concurrent Programs

int X = F = 0;

X = 1;  
F = 1;

t = F;  
u = X;

- What does this program mean?
- Sequential Consistency [Lamport '79]  
Program behavior = set of its thread interleavings

# Recall: Semantics of WHILE<sub>||</sub> from midterm

$$\frac{\langle E, S_1 \rangle \rightarrow \langle E', S'_1 \rangle}{\langle E, S_1; S_2 \rangle \rightarrow \langle E', S'_1; S_2 \rangle} \text{ small-seq-congruence}$$

$$\frac{}{\langle E, \text{skip}; S_2 \rangle \rightarrow \langle E, S_2 \rangle} \text{ small-seq}$$

$$\frac{\langle E, S_1 \rangle \rightarrow \langle E', S'_1 \rangle}{\langle E, S_1 \parallel S_2 \rangle \rightarrow \langle E', S'_1 \parallel S_2 \rangle} \text{ small-par-congruence-1}$$

$$\frac{\langle E, S_2 \rangle \rightarrow \langle E', S'_2 \rangle}{\langle E, S_1 \parallel S_2 \rangle \rightarrow \langle E', S_1 \parallel S'_2 \rangle} \text{ small-par-congruence-2}$$

$$\frac{}{\langle E, \text{skip} \parallel \text{skip} \rangle \rightarrow \langle E, \text{skip} \rangle} \text{ small-par-skip}$$

# Exercise 1:

```
int X = F = 0;
```

```
X = 1;  
F = 1;
```

```
t = F;  
u = X;
```

- What are the possible final values for variables `t` and `u` after running this program, assuming sequential consistency?

# Sequential Consistency Explained

int X = F = 0; // F = 1 implies X is initialized

X = 1;  
F = 1;

t = F;  
u = X;

X = 1;

X = 1;

X = 1;

t = F;

t = F;

t = F;

F = 1;

t = F;

t = F;

u = X;

X = 1;

X = 1;

t = F;

F = 1;

u = X;

X = 1;

u = X;

F = 1;

u = X;

u = X;

F = 1;

F = 1;

F = 1;

u = X;

t=1, u=1

t=0, u=1

t=0, u=1

t=0, u=0

t=0, u=1

t=0, u=1

t=1 implies u=1

# Naturalness of Sequential Consistency

- Sequential Consistency provides two crucial abstractions
- Program Order Abstraction
  - Instructions execute in the order specified in the program  
 $A ; B$   
means “Execute A and then B”
- Shared Memory Abstraction
  - Memory behaves as a global array, with reads and writes done immediately
- We implicitly assume these abstractions for sequential programs
  - As we will see, we can only rely on these abstractions under certain conditions in a concurrent context

# WHAT IS A DATA RACE ?

---

- The term “data race” is often overloaded to mean different things
- Precise definition is important in designing a tool

# Data Race

- Two accesses *conflict* if
  - they access the same memory location, and
  - at least one of them is a write

Write X – Write X

Write X – Read X


Read X – Write X

Read X – Read X

- A data race is a pair of conflicting accesses **that happen concurrently**

# “Happen Concurrently”

- A and B happen concurrently if they occur in different threads, and
- there exists a sequentially consistent execution in which they occur one after the other

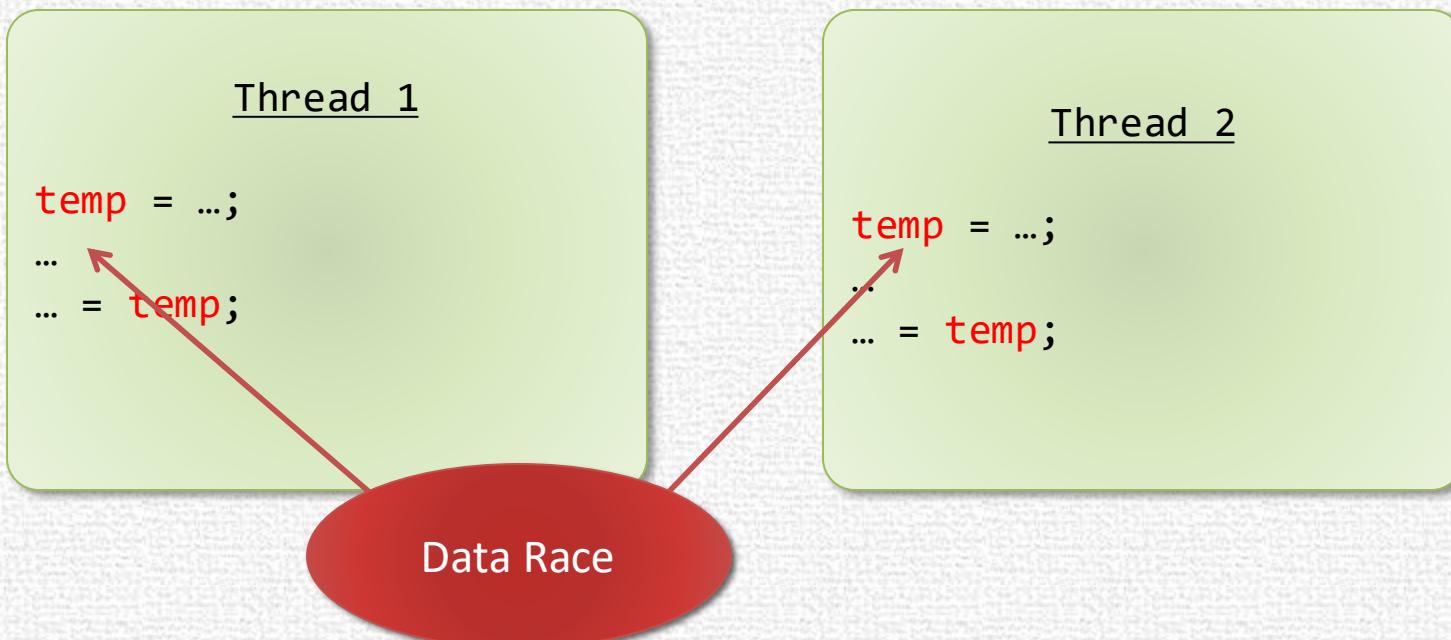


# Data races are almost always no good

- What are some consequences of a data race, even when assuming sequential consistency?

# Unintended Sharing

- Threads accidentally sharing data that should not be global
- *Solution:* Change allocation (e.g., stack var or static thread-local)



# Atomicity Violation

- When code that is meant to execute *atomically* (that is, perform a single undivisible operation) suffers interference from some other thread
- *Solution:* Surround critical sections with locks

Thread 1

```
void Bank::Update(int a)
{
    int t = bal;
    bal = t + a;
}
```


Thread 2

```
void Bank::Withdraw(int a)
{
    int t = bal;
    bal = t - a;
}
```

Data Race

# Ordering Violation

- Incorrect signaling between a producer and a consumer
- *Solution:* Reorder operations or use synchronization (e.g., signals)



# But,....

- How do you think "locks" are implemented?
- Atomic compare-and-swap (CAS)

```
AcquireLock(lock){  
    while (!CAS (lock, 0, 1)) {}  
}
```

```
ReleaseLock(lock) {  
    lock = 0;  
}
```

Data Race ?

# Acceptable Concurrent Conflicting Accesses

- Implementing synchronization (such as locks) usually requires concurrent conflicting accesses to shared memory
- Innovative uses of shared memory
  - Fast reads
  - Double-checked locking
  - Lazy initialization
  - Setting dirty flag
  - ...
- Need mechanisms to distinguish these from erroneous conflicts

# Solution: Programmer Annotation

- Programmer explicitly annotates variables as “synchronization”
  - Java – volatile keyword
  - C++ – std::atomic<> types

# Data Race

- Two accesses *conflict* if
  - they access the same memory location, and
  - at least one of them is a write
- A data race is a pair of concurrent conflicting accesses to locations **not annotated as synchronization**
  - Recall: “Concurrent” means there exists a sequentially consistent execution in which they happen one after the other
- Equivalent definition: a pair of conflicting accesses where one doesn’t **happen before** the other
  - Program order
  - Synchronization order
    - Acquire/release, wait-notify, fork-join, volatile read/write

# Exercise 2: Is there a data race?

## If so, on what variable(s)?

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

# Is there a data race?

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

# Consider regular compiler transformations/optimizations

Before:

```
data = 42;  
flag = true;
```

After:

```
flag = true;  
data = 42;
```

# Possible behavior

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

**data = 42;**

# Consider regular compiler transformations/optimizations

Before:

```
if (flag)  
    t = data;
```

After:

```
t2 = data;  
if (flag)  
    t = t2;
```

# Possible behavior

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
t2 = data;  
  
if (flag)  
    t = t2;
```

# How do we fix this?

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

# Using “synchronized” keyword in Java

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = ...;  
synchronized (m) {  
    flag = true;  
}
```

T2:

```
boolean f;  
synchronized (m) {  
    f = flag;  
}  
if (f)  
    ... = data;
```

# ... Implemented via locks

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = ...;  
acquire(m);  
flag = true;  
release(m);
```

T2:

```
boolean f;  
acquire(m);  
f = flag;  
release(m);  
if (f)  
    ... = data;
```

Happens-before  
relationship

# Using “volatile” keyword in Java

Initially:

```
int data = 0;  
volatile boolean flag = false;
```

T1:

```
data = ...;  
flag = true;
```

T2:

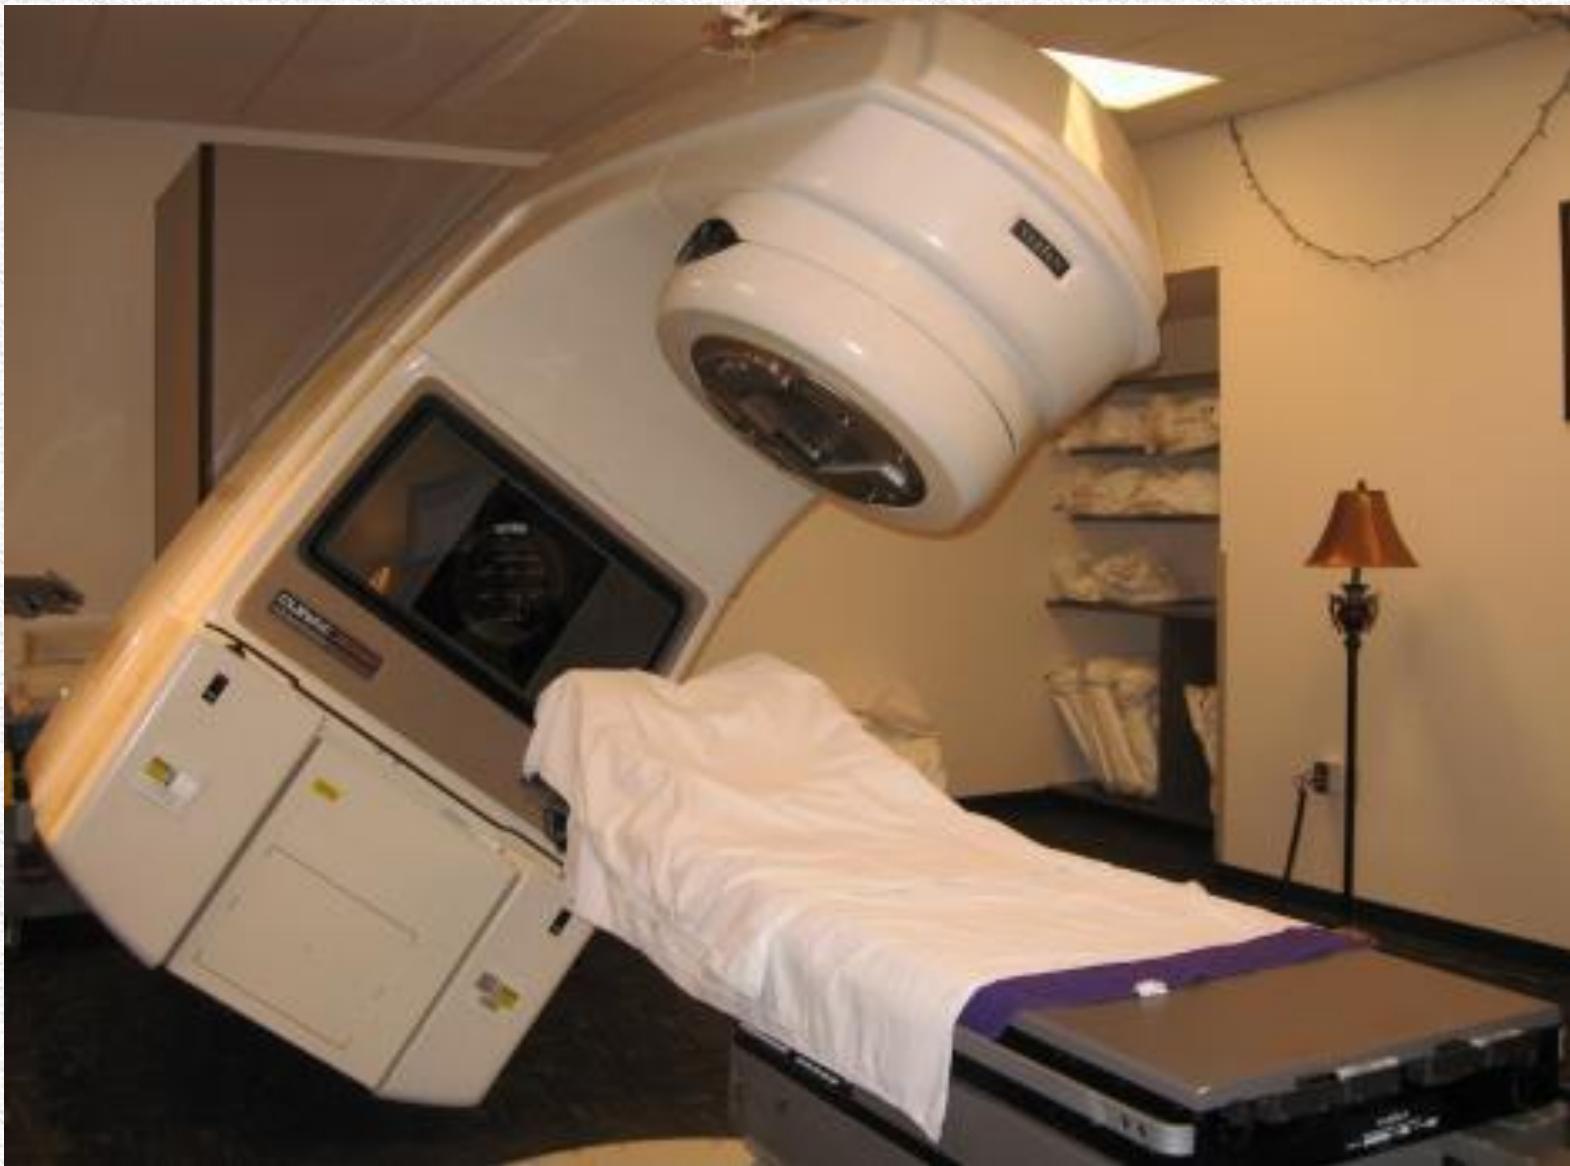
```
if (flag)  
    ... = data;
```

*Happens-before  
relationship*

# Data Race vs Race Conditions

- Data Races != Race Conditions
  - Confusing terminology
- Race Condition
  - Any timing error in the program
  - Due to events, device interaction, thread interleaving, ...
  - **Race conditions can be very bad!**





# Data Race vs Race Conditions

- Data Races != Race Conditions
  - Confusing terminology
- Race Condition
  - Any timing error in the program
  - Due to events, device interaction, thread interleaving, ...
  - **Race conditions can be very bad!**
- Data races are neither sufficient nor necessary for a race condition
  - Data race is a good **symptom** for a race condition


# DATA-RACE-FREEDOM SIMPLIFIES LANGUAGE SEMANTICS

---

# Advantage of Eliminating All Data Races

- Defining semantics for concurrent programs becomes surprisingly easy
- In the presence of compiler and hardware optimizations

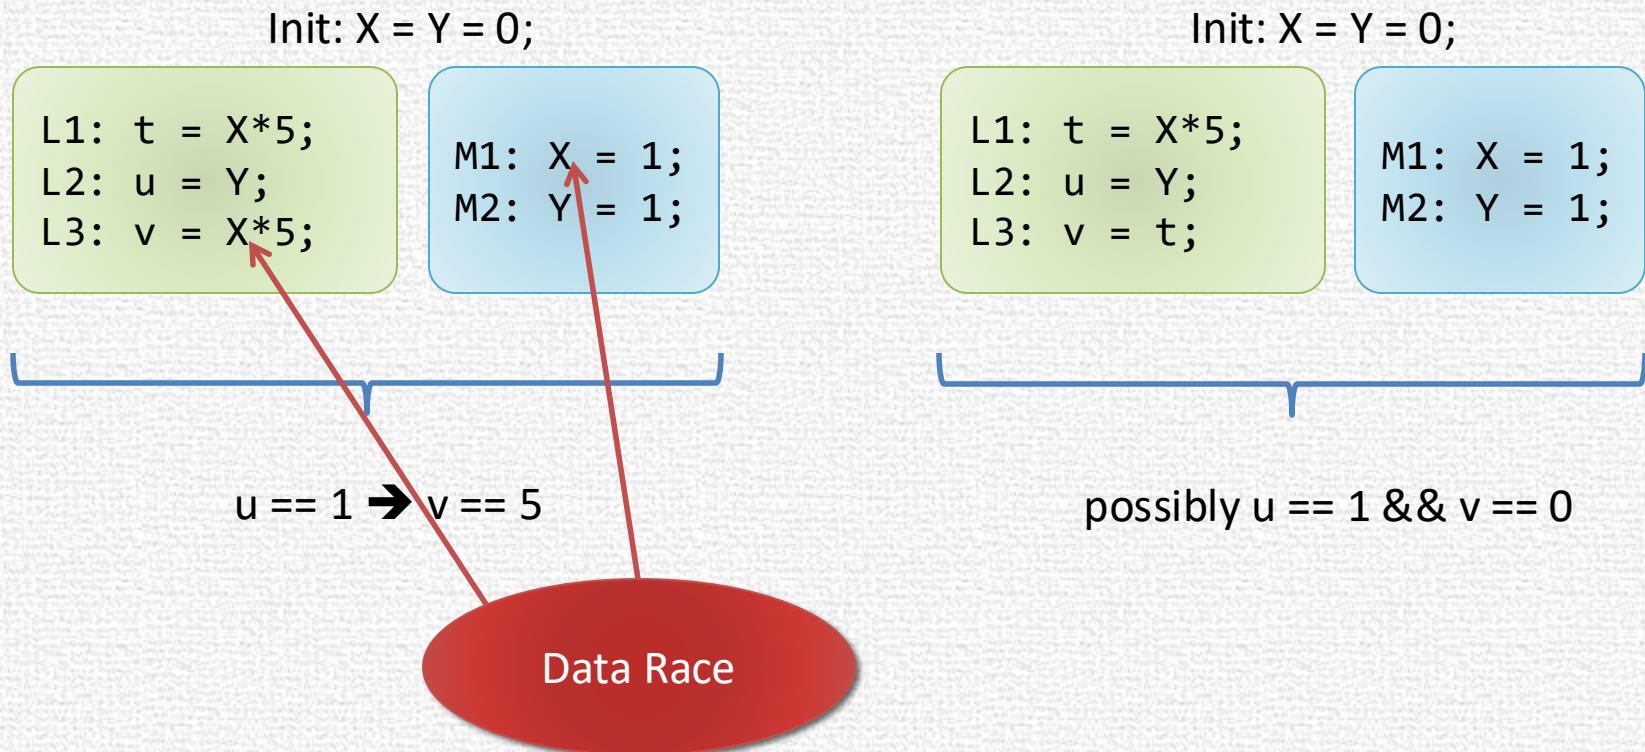
# Can A Compiler Do This?



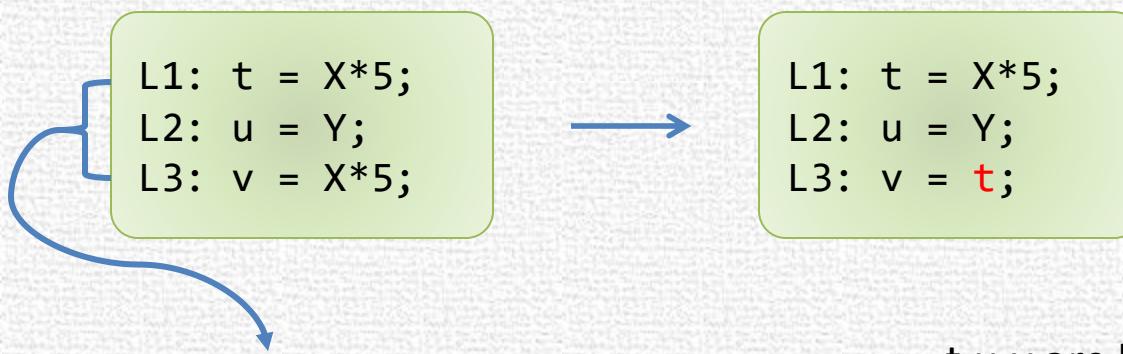
OK for sequential programs  
if X is not modified between L1 and L3

`t,u,v` are local variables  
`X,Y` are possibly shared

# Can Break Sequential Consistent Semantics



# Can A Compiler Do This?



OK for sequential programs  
if X is not modified between L1 and L3

t,u,v are local variables  
X,Y are possibly shared

OK for concurrent programs  
if there is no data race on X or  
if there is no data race on Y

# Key Observation [Adve& Hill '90 ]

- Many sequentially valid (compiler & hardware) transformations also preserve sequential consistency
- Provided the program is data-race free
- Forms the basis for modern C++, Java semantics
  - data-race-free → sequential consistency
  - otherwise → weak/undefined semantics

# DATA RACE DETECTION

---

# Overview of Data Race Detection Techniques

- Static data race detection
- Dynamic data race detection
  - Lock-set
  - Happen-before
  - DataCollider

# Static Data Race Detection

- Advantages:
  - Reason about all inputs/interleavings
  - No run-time overhead
  - Adapt well-understood static-analysis techniques
  - Annotations to document concurrency invariants
- Example Tools:
  - RCC/Java type-based
  - ESC/Java "functional verification"  
(theorem proving-based)

# Static Data Race Detection

- Advantages:
  - Reason about all inputs/interleavings
  - No run-time overhead
  - Adapt well-understood static-analysis techniques
  - Annotations to document concurrency invariants
- Disadvantages of static:
  - Undecidable...
  - Tools produce “false positives” or “false negatives”
  - May be slow, require programmer annotations
  - May be hard to interpret results

# Dynamic Data Race Detection

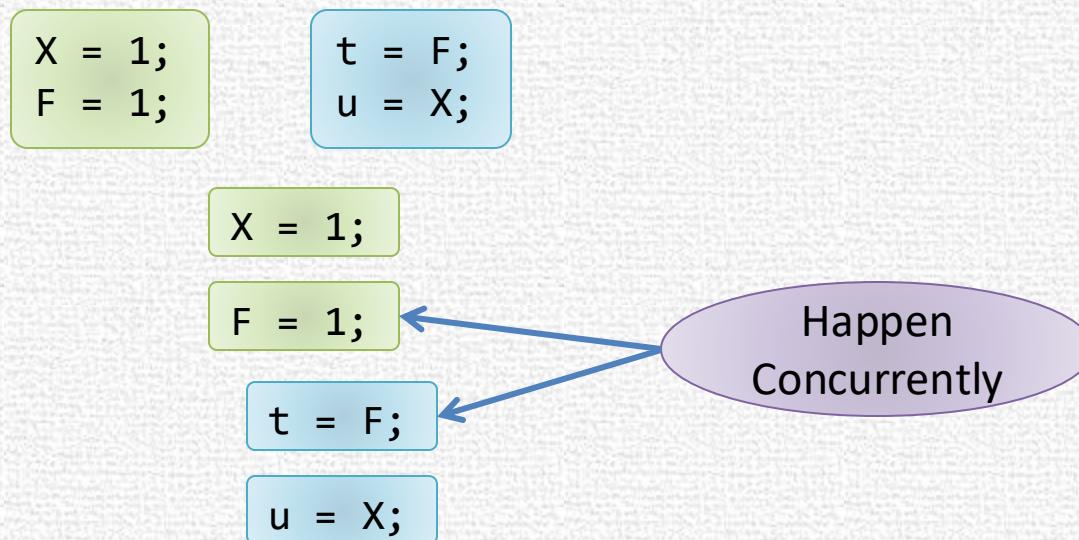
- Advantages
  - Can avoid “false positives”
  - No need for language extensions or sophisticated static analysis
- Disadvantages
  - Run-time overhead (5-20x for best tools)
  - Memory overhead for analysis state
  - Reasons only about observed executions
    - sensitive to test coverage
    - (some generalization possible...)

# Tradeoffs: Static vs Dynamic

- Coverage
  - generalize to additional traces?
- Soundness
  - all reported warnings are actually races
- Completeness
  - every actual data race is reported
- Overhead
  - run-time slowdown
  - memory footprint
- Programmer overhead

# Definition Refresh

- A data race is a pair of concurrent conflicting accesses to unannotated locations (i.e. not locks or volatile variables)



- Problem for dynamic data race detection
  - Very difficult to catch the two accesses executing concurrently

# Solution

- Lockset
  - Infer data races through violation of locking discipline
- Happens-before
  - Infer data races by generalizing a trace to a set of traces with the same happens-before relation

# LOCKSET ALGORITHM

---

Eraser [Savage et.al. '97]

# Lockset Algorithm Overview

- Checks a sufficient condition for data-race-freedom
- Consistent locking discipline
  - Every data structure is protected by a single lock
  - All accesses to the data structure made while holding the lock
- Example:

```
// Remove a received packet  
AcquireLock( RecvQueueLk );  
pkt = RecvQueue.RemoveTop();  
ReleaseLock( RecvQueueLk );
```

```
... // process pkt
```

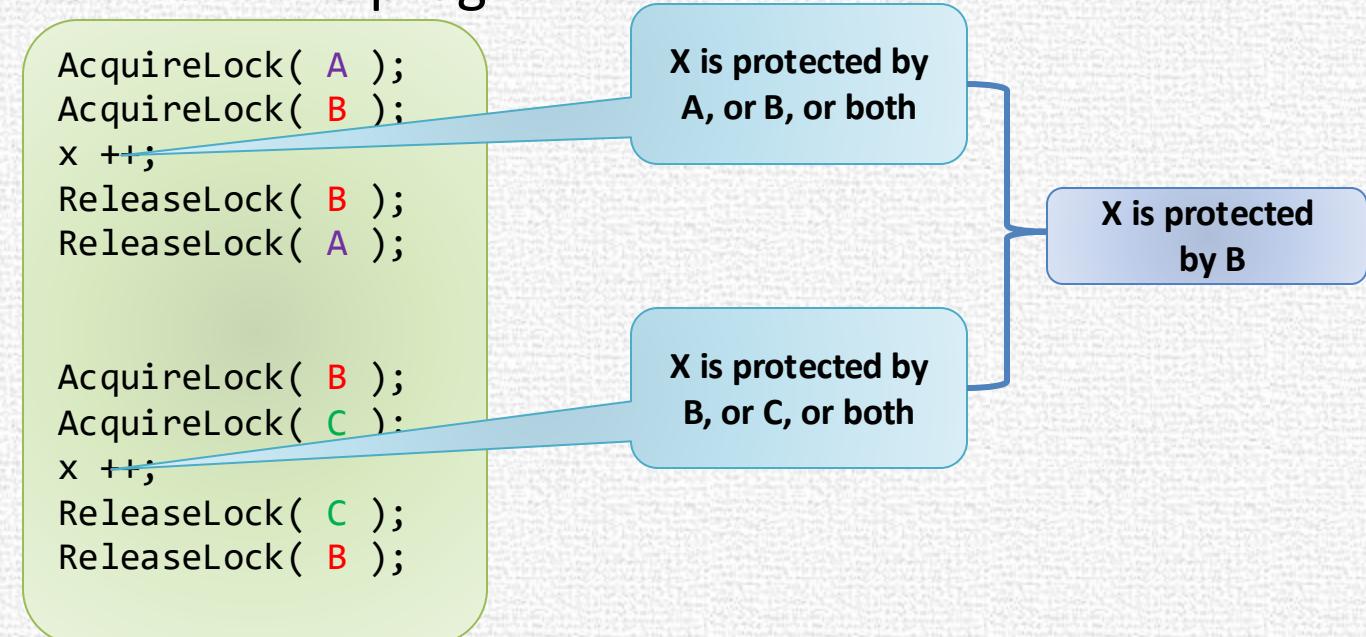
```
// Insert into processed  
AcquireLock( ProcQueueLk );  
ProcQueue.Insert(pkt);  
ReleaseLock( ProcQueueLk );
```

RecvQueue is  
consistently protected  
by RecvQueueLk

ProcQueue is  
consistently protected  
by ProcQueueLk

# Inferring the Locking Discipline

- How do we know which lock protects what?
  - Asking the programmer is cumbersome
- Solution: Infer from the program



# LockSet Algorithm

- Two data structures:
  - $\text{LocksHeld}( t )$  = set of locks held currently by thread  $t$ 
    - Initially set to Empty
  - $\text{LockSet}( x )$  = set of locks that could potentially be protecting  $x$ 
    - Initially set to the universal set
- When thread  $t$  acquires lock  $l$ 
  - $\text{LocksHeld}( t ) = \text{LocksHeld}( t ) \cup \{l\}$
- When thread  $t$  releases lock  $l$ 
  - $\text{LocksHeld}( t ) = \text{LocksHeld}( t ) - \{l\}$
- When thread  $t$  accesses location  $x$ 
  - $\text{LockSet}( x ) = \text{LockSet}( x ) \cap \text{LocksHeld}( t )$
  - Report “data race” when  $\text{LockSet}( x )$  becomes empty

# LockSet Algorithm

- No warnings → no data races on the current execution
  - The program followed consistent locking discipline in this execution
- Warnings does not imply a data race
  - Thread-local initialization

```
// Initialize a packet  
pkt = new Packet();  
pkt.Consumed = 0  
  
AcquireLock( SendQueueLk );  
SendQueue.Enqueue(pkt);  
ReleaseLock( SendQueueLk );
```

```
// Process a packet  
AcquireLock( SendQueueLk );  
pkt = SendQueue.Top();  
pkt.Consumed = 1;  
ReleaseLock( SendQueueLk );
```

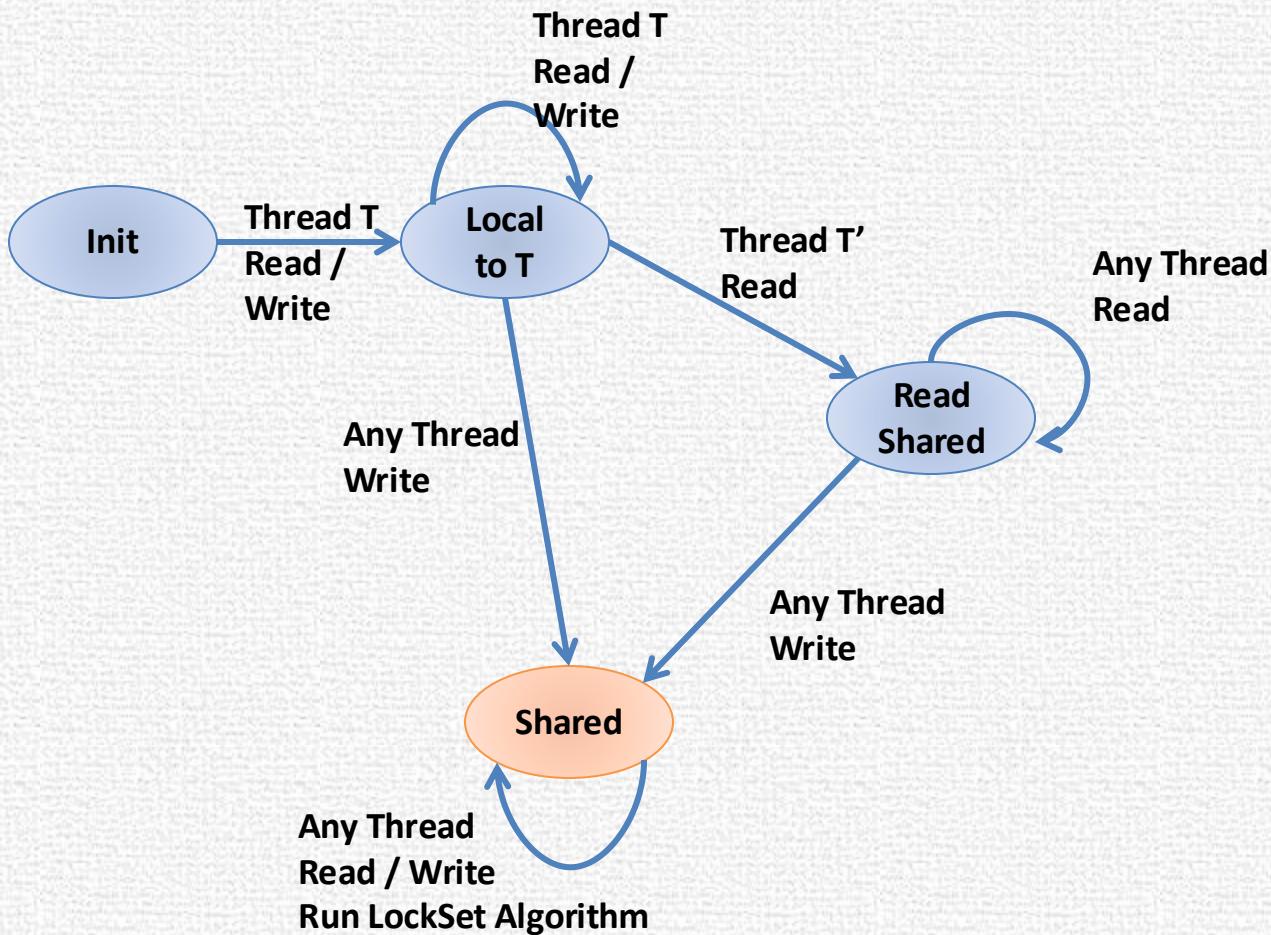
# LockSet Algorithm

- No warnings → no data races on the current execution
  - The program followed consistent locking discipline in this execution
- Warnings does not imply a data race
  - Object read-shared after thread-local initialization

```
A = new A();  
A.f = 0;  
  
// publish A  
globalA = A;
```

```
f = globalA.f;
```

# Maintain A State Machine Per Location



# LockSet Algorithm

- State machine misses some data races

```
// Initialize a packet  
pkt = new Packet();  
pkt.Consumed = 0;
```

```
AcquireLock( WrongLk );  
pkt = SendQueue.Top();  
pkt.Consumed = 1;  
ReleaseLock( WrongLk );
```

```
// Process a packet  
AcquireLock( SendQueueLk );  
pkt = SendQueue.Top();  
pkt.Consumed = 1;  
ReleaseLock( SendQueueLk );
```

# LockSet Algorithm

- Does not handle locations consistently protected by different locks during a particular execution

```
// Remove a received packet  
AcquireLock( RecvQueueLk );  
pkt = RecvQueue.RemoveTop();  
ReleaseLock( RecvQueueLk );
```

```
... // process pkt
```

```
// Insert into processed  
AcquireLock( ProcQueueLk );  
ProcQueue.Insert(pkt);  
ReleaseLock( ProcQueueLk );
```

Pkt is protected by  
RecvQueueLk

Pkt is thread local


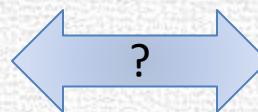
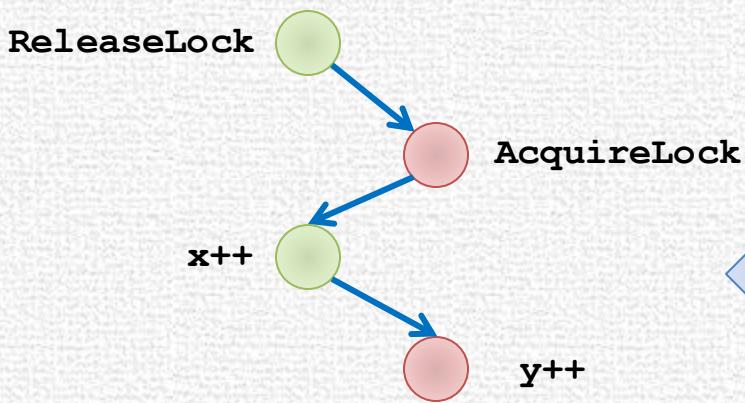
Pkt is protected by  
ProcQueueLk

HAPPENS-BEFORE

---

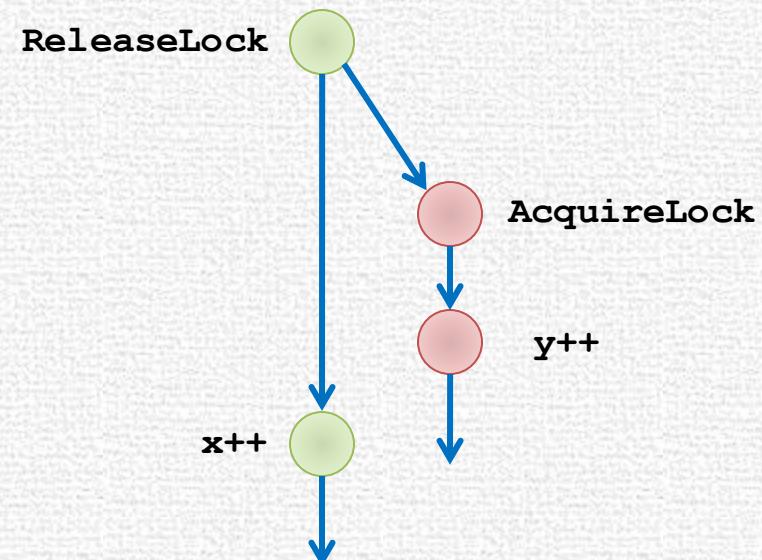

# Happens-Before Relation [Lamport '78]

- A concurrent execution is a partial-order determined by communication events
- The program cannot “observe” the order of concurrent non-communicating events



# Happens-Before Relation [Lamport '78]

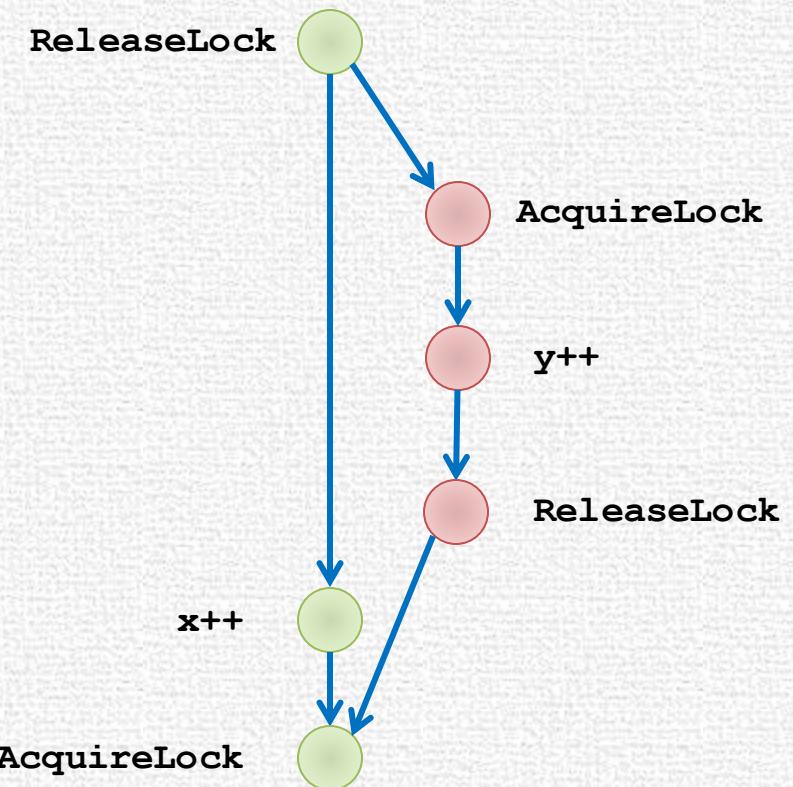
- A concurrent execution is a partial-order determined by communication events
- The program cannot “observe” the order of concurrent non-communicating events



- Both executions form the same happens-before relation


# Constructing the Happens-Before Relation

- Program order
  - Total order of thread instructions
- Synchronization order
  - Total order of accesses to the same synchronization



# Happens-Before Relation And Data Races


- If all conflicting accesses are ordered by happens-before
  - → data-race-free execution
  - → All linearizations of partial-order are valid program executions
- 
- If there exists conflicting accesses not ordered
  - → a data race




# IMPLEMENTING HAPPENS-BEFORE ANALYSES


---


# Dynamic Data-Race Detection





# Precise Happens-Before














# Exercise on vector clocks and partial ordering

- $VC = [t_1, t_2, \dots, t_N]$
- What is  $VC_a \sqsubseteq VC_b$ ?
- What is  $VC_a \sqcup VC_b$ ?
- What are sufficient and necessary conditions for there to be a data race between two accesses having vector clocks  $VC_a$  and  $VC_b$ ?

$VC_A$ 

4	1
---	---

A B

 $VC_B$ 

2	8
---	---

A B

 $L_m$ 

2	1
---	---

A B

 $W_x$ 

3	0
---	---

A B

 $R_x$ 

0	1
---	---

A B

A's local time

B's local time

$VC_A$ 

4	1
---	---

A      B

 $VC_B$ 

2	8
---	---

A      B

 $L_m$ 

2	1
---	---

A      B

 $W_x$ 

3	0
---	---

A      B


 $R_x$ 

0	1
---	---

A      B

B-steps with B-time  $\leq 1$   
happen before  
A's next step





Write-Write Check:  $W_x \sqsubseteq VC_A$  ?

3	0
---	---

 $\sqsubseteq$ 

4	1
---	---

 ? Yes

Read-Write Check:  $R_x \sqsubseteq VC_A$  ?


0	1
---	---

 $\sqsubseteq$ 

4	1
---	---

 ? Yes

$O(n)$  time



$VC_A$ 

4	1
---	---

$x = 0$

4	1
---	---

 $rel(m)$ 

5	1
---	---

 $VC_B$ 

2	8
---	---

 $L_m$ 

2	1
---	---

 $W_x$ 

3	0
---	---

 $R_x$ 

0	1
---	---


2	8
---	---


2	8
---	---


2	1
---	---


4	0
---	---

0	1
---	---










# VectorClocks for Data-Race Detection

- Sound
  - Warning → data-race exists
- Complete
  - No warnings → data-race-free execution
- Performance
  - slowdowns > 50x
  - memory overhead


# Dynamic Data-Race Detection




# FASTTRACK


---

# Dynamic Data-Race Detection



# Dynamic Data-Race Detection





Write-Write Check:  $W_x \sqsubseteq VC_A$  ?

3	0
---	---

 $\sqsubseteq$ 

4	1
---	---

 ? Yes

Read-Write Check:  $R_x \sqsubseteq VC_A$  ?

0	1
---	---

 $\sqsubseteq$ 


4	1
---	---

 ? Yes

$O(n)$  time

# Write-Write and Write-Read Data Races

Thread A



$x = 0$

Thread B



$x = 1$

?


Thread C



$x = 4$


?

Thread D




$x = 3$

$O(n)$




# No Data Races Yet: Writes Totally Ordered

Thread A      Thread B      Thread C      Thread D



# No Data Races Yet: Writes Totally Ordered


Thread A      Thread B      Thread C      Thread D



$VC_A$

4	1
---	---

$x = 0$



$VC_B$

2	8
---	---

$L_m$

2	1
---	---

$W_x$

1@B
-----


Last Write  
"Epoch"


Write-Write Check:  $W_x \sqsubset VC_A$  ?

$1@B \sqsubset [4 \ 1] ?$  Yes


$(1 \leq 1?)$

$O(1)$  time






# Read-Write Data Races -- Ordered Reads



Most common case: thread-local, lock-protected, ...

# Read-Write Data Races -- Unordered Reads



$VC_A$

7	0
---	---

$x = 0$

7	0
---	---

**fork**

8	0
---	---

7	1
---	---

**read x**

8	0
---	---

**read x**

8	0
---	---

$x = 2$

--	--

$VC_B$

7	1
---	---

7	1
---	---

$W_x$

-
---

-
---

7@A
-----

-
---

7@A
-----

-
---

7@A
-----

1@B
-----

7@A
-----

8 1
-----

$O(1)$

$O(n)$

$O(n)$

Read-Write Check:  $R_x \subseteq VC_A$  ?

8	1
---	---

8	0
---	---

?


No

Thread A


Thread B

Thread C


Thread D



Thread A




Thread B



Thread C




Thread D



**read  $y$**

**read  $y$**

**$y = 10$**




Thread A


Thread B

Thread C


Thread D



### Thread A



### Thread B




### Thread C



### Thread D



Forget VC for  $R_x$   
and switch back  
to "last read epoch"




$y = 10$

?

$y = 3$

$O(1)$

# Slowdown (x Base Time)



# Memory Usage

- FastTrack allocated ~200x fewer VCs


Checker	Memory Overhead
Basic VC, DJIT+	7.9x
FastTrack	2.8x
Empty	2.0x

(Note: VCs for dead objects are garbage collected)

- Improvements
  - accordion clocks [CB 01]
  - analysis granularity [PS 03, YRC 05]

# Eclipse 3.4

- Scale
  - > 6,000 classes
  - 24 threads
  - custom sync. idioms
- Precision (tested 5 common tasks)
  - Eraser: ~1000 warnings
  - FastTrack: ~30 warnings
- Performance on compute-bound tasks
  - > 2x speed of other precise checkers
  - same as Eraser



# FUZZING TECHNIQUES

---

# Fuzzing can also find data races

- Idea: Catch races “red handed”. Loosely,
  - Pause thread execution when writing to X
  - If another thread reaches a statement that reads or writes X then we have observed concurrent conflicting accesses!
- Analysis does not care about locks or other synchronization primitives.
  - Consistent locking will make the above condition impossible to trigger.

# Race Fuzzer

- Run-time Overhead
  - No overhead of tracking synchronization, locks, or vector clocks (hey, that rhymes!)
  - But pausing threads forever can lead to deadlocks
  - Pausing threads for a short while (e.g. `sleep(1000)`) adds overhead for every write access, though this approach is very effective.
- Solution idea:
  - Instead of “pausing” thread, just deprioritize it in the OS scheduler

# Race Fuzzing

- Randomized scheduling still depends on luck
- Can do systematic schedule exploration with a bounded number of context switches
- Sophisticated randomized algorithms like PCT can give probabilistic guarantees of uncovering concurrency bugs with a bounded number of “ordering constraints”.
- Or use heuristics, e.g. TSVD uses an initial run to infer “likely” happens-before relationships based on wall-clock timestamps to select candidate “racing pairs”.

# Lecture Takeaways

- Data race: two accesses, one of which is a write, with no happens-before relation
- Data races are subtle
  - Compiler optimizations, hardware reordering make racy program behavior hard to predict
  - Better to synchronize consistently
- Lockset analysis: intuitive, fast
  - But many false warnings
- Happens-before data race detection
  - Sound; OK speed if carefully implemented
- Stress testing
  - Sound and fast; Can catch data races red handed
  - Needs assumptions to prune the space of possible races

# Key References

- Hans-J. Boehm and Sarita V. Adve, "You Don't Know Jack About Shared Variables or Memory Models", CACM 2012.
- Leslie Lamport, "Time, Clocks, and the Ordering of Events in a Distributed System", CACM 1978.
- Martin Abadi, Cormac Flanagan, and Stephen N. Freund, "Types for Safe Locking: Static Race Detection for Java", TOPLAS 2006.
- Madanlal Musuvathi, Shaz Qadeer, Thomas Ball, Gerard Basler, Piramanayagam Arumuga Nainar, and Iulian Neamtiu, "Finding and Reproducing Heisenbugs in Concurrent Programs", OSDI 2008.
- Cormac Flanagan, K. Rustan M. Leino, Mark Lillibridge, Greg Nelson, James B. Saxe, and Raymie Stata. "Extended static checking for Java", PLDI 2002.
- S. Savage, M. Burrows, G. Nelson, P. Sobalvarro, and T. E. Anderson, "Eraser: A dynamic data race detector for multi-threaded programs", TOCS 1997.

# Key References

- Friedemann Mattern, "Virtual Time and Global States of Distributed Systems", Workshop on Parallel and Distributed Algorithms 1989.
- Yuan Yu, Tom Rodeheffer, and Wei Chen, "RaceTrack: Efficient detection of data race conditions via adaptive tracking", SOSP 2005.
- Eli Pozniansky and Assaf Schuster, "MultiRace: Efficient on-the-fly data race detection in multithreaded C++ programs", Concurrency and Computation: Practice and Experience 2007.
- Robert O'Callahan and Jong-Deok Choi, "Hybrid Dynamic Data Race Detection", PPOPP 2003.
- Cormac Flanagan and Stephen N. Freund, "FastTrack: efficient and precise dynamic race detection", CACM 2010.
- Cormac Flanagan and Stephen N. Freund, "The RoadRunner dynamic analysis framework for concurrent programs", PASTE 2010.

# Key References

- John Erickson, Madanlal Musuvathi, Sebastian Burckhardt, Kirk Olynyk, "Effective Data-Race Detection for the Kernel", OSDI 2010.
- Madanlal Musuvathi, Sebastian Burckhardt, Pravesh Kothari, and Santosh Nagarakatte, "A Randomized Scheduler with Probabilistic Guarantees of Finding Bugs", ASPLOS 2010.
- Michael D. Bond, Katherine E. Coons, Kathryn S. McKinley, "PACER: proportional detection of data races", PLDI 2010.
- Cormac Flanagan and Stephen N. Freund, "Adversarial memory for detecting destructive races", PLDI 2010.
- Koushik Sen. "Race directed random testing of concurrent programs". PLDI 2010.
- Guangpu Li, Shan Lu, Madanlal Musuvathi, Suman Nath, and Rohan Padhye. "Efficient scalable thread-safety-violation detection: finding thousands of concurrency bugs during testing", SOSP 2019.

# Bonus slides on the Java Memory Model (JMM)

# Behaviors Allowed in JMM

```
int data = flag = 0;
```

T1


```
r = data;  
flag = 1;
```

```
assert r == 0;
```


T2

```
while (flag == 0) {}  
data = 1;
```


# Behaviors Allowed in JMM




# Behaviors Allowed in JMM



# Behaviors Allowed in JMM



# Behaviors Allowed in JMM



# Behaviors Allowed in JMM

```
int data = flag = 0;
```

T1

```
r = data;  
flag = 1;
```

T2

```
while (flag == 0) {}  
data = 1;
```

```
assert r == 0;
```

Assertion  
failure!

# Behaviors Allowed in JMM

```
int data = flag = 0;
```

T1

```
r = data;  
flag = 1;  
assert r == 0;
```

T2

```
while (flag == 0) {}  
data = 1;
```

Requires returning future value or  
reordering to trigger the assertion failure

# Can this assert trigger in JVMs? Do you think the JMM allows it?

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = y;  
    x = r3;  
} else x = 1;
```

```
assert r2 == 0;
```

# The JVM and the JMM

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = y;  
    x = r3;  
} else x = 1;
```

JMM disallows  
r2 == 1 because  
of causality  
requirements

```
assert r2 == 0;
```

– Ševčík and Aspinall, ECOOP, 2008

# The JVM and the JMM

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

However, in a  
**JVM**, after  
redundant read  
elimination

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = r2;  
    x = r3;  
} else x = 1;
```

```
assert r2 == 0;
```

# The JVM and the JME

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

However, in a  
**JVM**, after  
redundant read  
elimination

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = r2;  
    x = r3;  
} else x = 1;
```

```
assert r2 == 0;
```

# The JVM and the JME

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

However, in a  
**JVM**, after  
redundant read  
elimination

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = r2;  
    x = r3;  
} else x = 1;  
  
assert r2 == 0;      r2 = y;  
                      If (r2 == 1)  
                      x = r2;  
                      else x = 1;  
  
                      ↓  
                      r2 = y;  
                      x = 1;
```

# The JVM and the JME

```
int x = y = 0;
```

T1

```
r1 = x;  
y = r1;
```

However, in a  
**JVM**, after  
redundant read  
elimination

T2

```
r2 = y;  
if (r2 == 1) {  
    r3 = r2;  
    x = r3;  
} else x = 1;
```

```
r2 = y;  
If (r2 == 1)  
x = r2;  
else x = 1;
```



```
r2 = y;  
x = 1;
```

```
assert r2 == 0;
```

Assertion  
failure  
possible!

# Moral: Just say no to data races

Don't try hacks based on the memory model

- Unless you are as good as Doug Lea



Author of `java.util.concurrent`

- Or you have formalized the memory model rules in a tool
  - And even then, are the rules right?