

# CONCURRENCY: SEQUENTIAL CONSISTENCY, DATA RACES, AND DYNAMIC ANALYSES

---

Lecture by Rohan Padhye

17-355/17-665/17-819: Program Analysis

Material from past lectures by Jonathan Aldrich, based in large part on slides by John Erickson, Stephen Freund, Madan Musuvathi, Mike Bond, and Man Cao

# Lecture Goals

- What is sequential consistency and why is it important?
- What is a data race, and what is data-race-free execution?
- Subtleties of data races and memory models
  - Why taking advantage of “harmless races” is almost certainly a bad idea
- Lockset analysis for data race detection
- Happens-before based data race detection
  - And high performance implementations, e.g. as in FastTrack

# SEQUENTIAL CONSISTENCY

---

# First things First

## Assigning Semantics to Concurrent Programs

int X = F = 0;

X = 1;  
F = 1;

t = F;  
u = X;

- What does this program mean?
- Sequential Consistency [Lamport '79]  
Program behavior = set of its thread interleavings

# Recall: Semantics of WHILE<sub>||</sub> from midterm

$$\frac{\langle E, S_1 \rangle \rightarrow \langle E', S'_1 \rangle}{\langle E, S_1; S_2 \rangle \rightarrow \langle E', S'_1; S_2 \rangle} \text{ small-seq-congruence}$$

$$\frac{}{\langle E, \text{skip}; S_2 \rangle \rightarrow \langle E, S_2 \rangle} \text{ small-seq}$$

$$\frac{\langle E, S_1 \rangle \rightarrow \langle E', S'_1 \rangle}{\langle E, S_1 \parallel S_2 \rangle \rightarrow \langle E', S'_1 \parallel S_2 \rangle} \text{ small-par-congruence-1}$$

$$\frac{\langle E, S_2 \rangle \rightarrow \langle E', S'_2 \rangle}{\langle E, S_1 \parallel S_2 \rangle \rightarrow \langle E', S_1 \parallel S'_2 \rangle} \text{ small-par-congruence-2}$$

$$\frac{}{\langle E, \text{skip} \parallel \text{skip} \rangle \rightarrow \langle E, \text{skip} \rangle} \text{ small-par-skip}$$

# Exercise 1:

```
int X = F = 0;
```

```
X = 1;  
F = 1;
```

```
t = F;  
u = X;
```

- What are the possible final values for variables `t` and `u` after running this program, assuming sequential consistency?

# Sequential Consistency Explained

int X = F = 0; // F = 1 implies X is initialized

X = 1;  
F = 1;

t = F;  
u = X;

X = 1;

X = 1;

X = 1;

t = F;

t = F;

t = F;

F = 1;

t = F;

t = F;

u = X;

X = 1;

X = 1;

t = F;

F = 1;

u = X;

X = 1;

u = X;

F = 1;

u = X;

u = X;

F = 1;

F = 1;

F = 1;

u = X;

t=1, u=1

t=0, u=1

t=0, u=1

t=0, u=0

t=0, u=1

t=0, u=1

t=1 implies u=1

# Naturalness of Sequential Consistency

- Sequential Consistency provides two crucial abstractions
- Program Order Abstraction
  - Instructions execute in the order specified in the program  
 $A ; B$   
means “Execute A and then B”
- Shared Memory Abstraction
  - Memory behaves as a global array, with reads and writes done immediately
- We implicitly assume these abstractions for sequential programs
  - As we will see, we can only rely on these abstractions under certain conditions in a concurrent context

# WHAT IS A DATA RACE ?

---

- The term “data race” is often overloaded to mean different things
- Precise definition is important in designing a tool

# Data Race

- Two accesses *conflict* if
  - they access the same memory location, and
  - at least one of them is a write

Write X – Write X

Write X – Read X


Read X – Write X

Read X – Read X

- A data race is a pair of conflicting accesses **that happen concurrently**

# “Happen Concurrently”

- A and B happen concurrently if
- there exists a sequentially consistent execution in which they happen one after the other

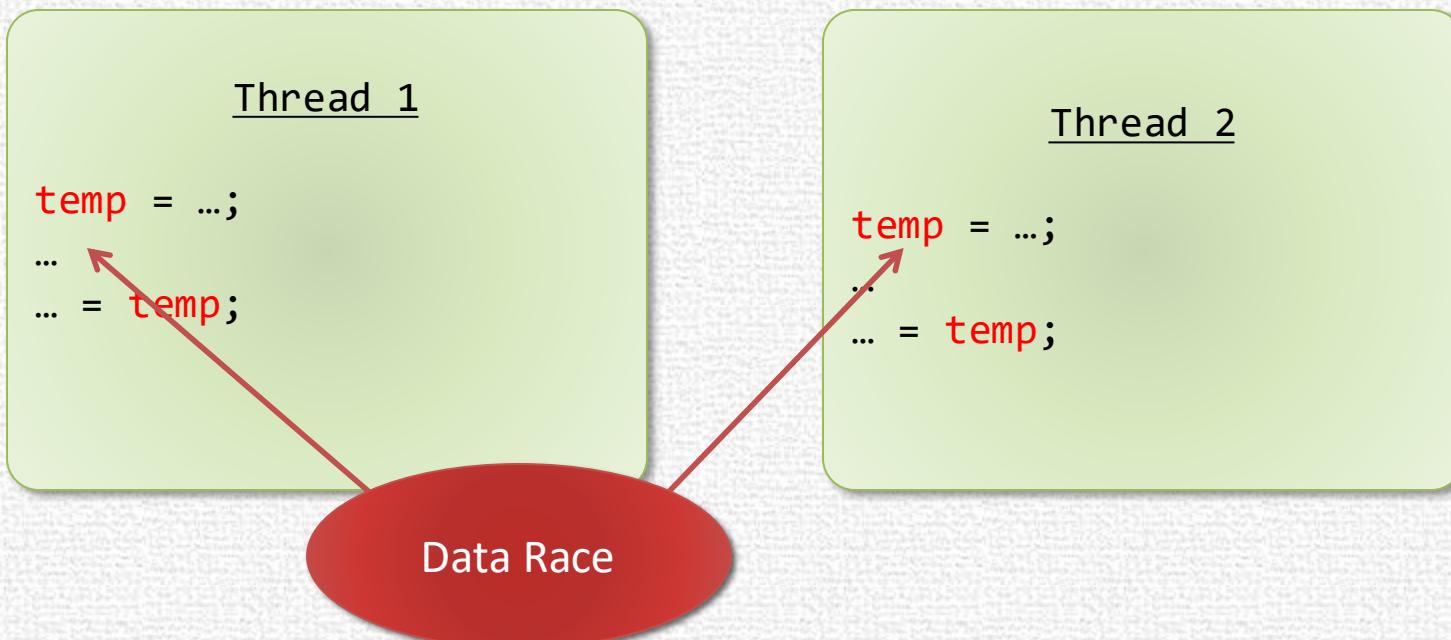


# Data races are almost always no good

- What are some consequences of a data race, even when assuming sequential consistency?

# Unintended Sharing

- Threads accidentally sharing data that should not be global
- *Solution:* Change allocation (e.g., stack var or static thread-local)



# Atomicity Violation

- When code that is meant to execute *atomically* (that is, perform a single undivisible operation) suffers interference from some other thread
- *Solution:* Surround critical sections with locks

Thread 1

```
void Bank::Update(int a)
{
    int t = bal;
    bal = t + a;
}
```


Thread 2

```
void Bank::Withdraw(int a)
{
    int t = bal;
    bal = t - a;
}
```

Data Race

# Ordering Violation

- Incorrect signaling between a producer and a consumer
- *Solution:* Reorder operations or use synchronization (e.g., signals)



# But,....

- How do you think "locks" are implemented?
- Atomic compare-and-swap (CAS)

```
AcquireLock(lock){  
    while (!CAS (lock, 0, 1)) {}  
}
```

```
ReleaseLock(lock) {  
    lock = 0;  
}
```

Data Race ?

# Acceptable Concurrent Conflicting Accesses

- Implementing synchronization (such as locks) usually requires concurrent conflicting accesses to shared memory
- Innovative uses of shared memory
  - Fast reads
  - Double-checked locking
  - Lazy initialization
  - Setting dirty flag
  - ...
- Need mechanisms to distinguish these from erroneous conflicts

# Solution: Programmer Annotation

- Programmer explicitly annotates variables as “synchronization”
  - Java – volatile keyword
  - C++ – std::atomic<> types

# Data Race

- Two accesses *conflict* if
  - they access the same memory location, and
  - at least one of them is a write
- A data race is a pair of concurrent conflicting accesses to locations **not annotated as synchronization**
  - Recall: “Concurrent” means there exists a sequentially consistent execution in which they happen one after the other
- Equivalent definition: a pair of conflicting accesses where one doesn’t **happen before** the other
  - Program order
  - Synchronization order
    - Acquire/release, wait-notify, fork-join, volatile read/write

# Exercise 2: Is there a data race?

## If so, on what variable(s)?

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

# Is there a data race?

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

# Consider regular compiler transformations/optimizations

Before:

```
data = 42;  
flag = true;
```

After:

```
flag = true;  
data = 42;
```

# Possible behavior

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

**data = 42;**

# Consider regular compiler transformations/optimizations

Before:

```
if (flag)  
    t = data;
```

After:

```
t2 = data;  
if (flag)  
    t = t2;
```

# Possible behavior

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
t2 = data;  
  
if (flag)  
    t = t2;
```

# How do we fix this?

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = 42;  
flag = true;
```

T2:

```
if (flag)  
    t = data;
```

# Using “synchronized” keyword in Java

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = ...;  
synchronized (m) {  
    flag = true;  
}
```

T2:

```
boolean f;  
synchronized (m) {  
    f = flag;  
}  
if (f)  
    ... = data;
```

# ... Implemented via locks

Initially:

```
int data = 0;  
boolean flag = false;
```

T1:

```
data = ...;  
acquire(m);  
flag = true;  
release(m);
```

T2:

```
boolean f;  
acquire(m);  
f = flag;  
release(m);  
if (f)  
    ... = data;
```

Happens-before  
relationship

# Using “volatile” keyword in Java

Initially:

```
int data = 0;  
volatile boolean flag = false;
```

T1:

```
data = ...;  
flag = true;
```

T2:

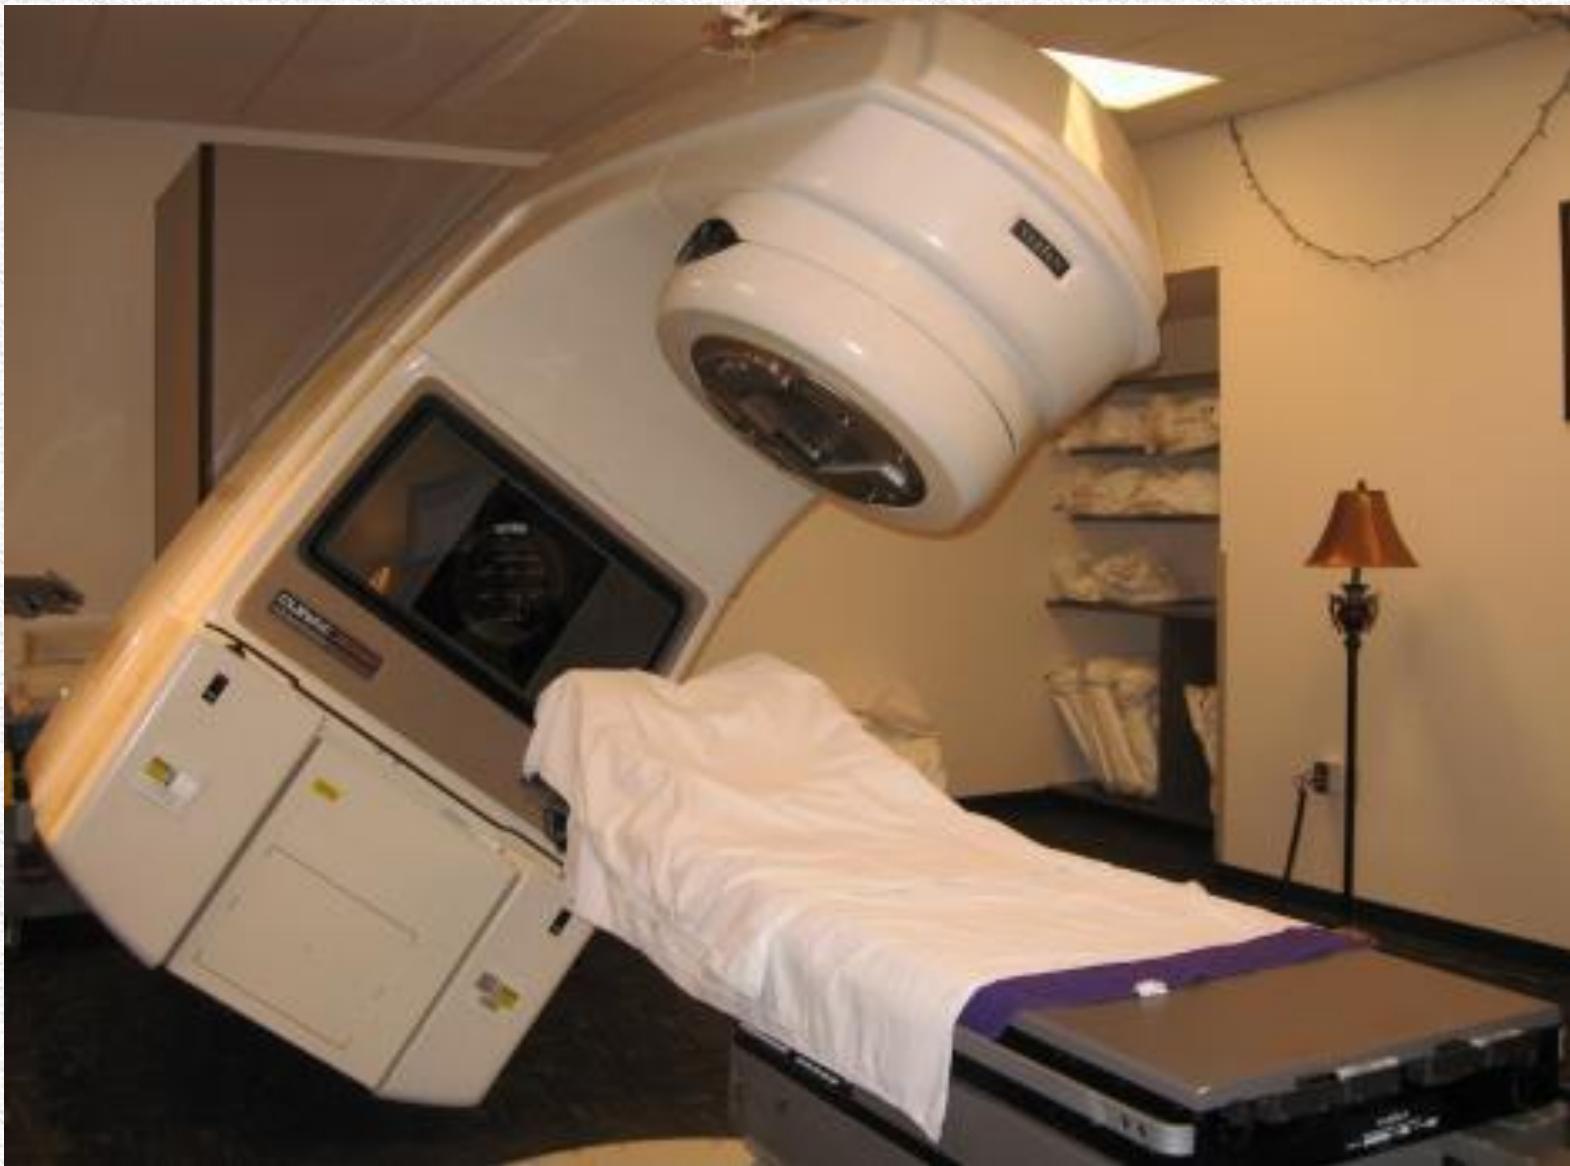
```
if (flag)  
    ... = data;
```

*Happens-before  
relationship*

# Data Race vs Race Conditions

- Data Races != Race Conditions
  - Confusing terminology
- Race Condition
  - Any timing error in the program
  - Due to events, device interaction, thread interleaving, ...
  - **Race conditions can be very bad!**





# Data Race vs Race Conditions

- Data Races != Race Conditions
  - Confusing terminology
- Race Condition
  - Any timing error in the program
  - Due to events, device interaction, thread interleaving, ...
  - **Race conditions can be very bad!**
- Data races are neither sufficient nor necessary for a race condition
  - Data race is a good **symptom** for a race condition


# DATA-RACE-FREEDOM SIMPLIFIES LANGUAGE SEMANTICS

---

# Advantage of Eliminating All Data Races

- Defining semantics for concurrent programs becomes surprisingly easy
- In the presence of compiler and hardware optimizations

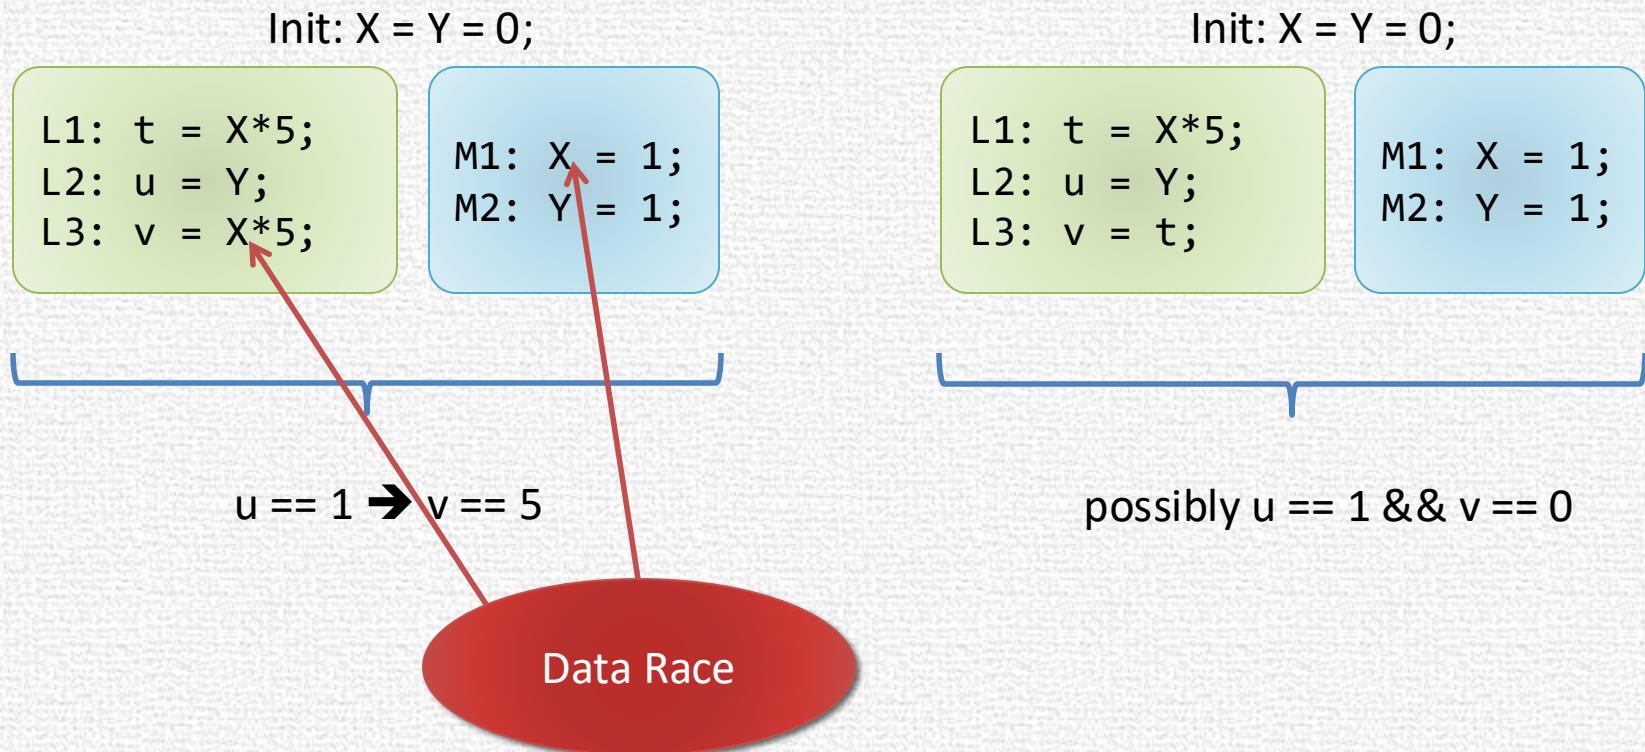
# Can A Compiler Do This?




OK for sequential programs  
if  $X$  is not modified between L1 and L3

$t, u, v$  are local variables  
 $X, Y$  are possibly shared

# Can Break Sequential Consistent Semantics



# Can A Compiler Do This?



OK for sequential programs  
if X is not modified between L1 and L3

t,u,v are local variables  
X,Y are possibly shared

OK for concurrent programs  
if there is no data race on X or  
if there is no data race on Y

# Key Observation [Adve& Hill '90 ]

- Many sequentially valid (compiler & hardware) transformations also preserve sequential consistency
- Provided the program is data-race free
- Forms the basis for modern C++, Java semantics
  - data-race-free → sequential consistency
  - otherwise → weak/undefined semantics